


Monitoring Kafka



Alex Loddengaard

January 25, 2019 18:08

[Follow](#)

System

Monitoring the following system metrics:

- CPU usage
- Memory usage
 - Note that the Kafka broker will only use JVM heap space for meta data and buffers. The broker relies on the OS' page cache for caching.
- Available disk space
- Disk IO
- Network IO
- Number of open file handles

Alert when the following is observed:

- 60% disk usage for disks storing the Kafka log (configured via log.dirs)
- 60% disk IO usage
- 60% network IO usage
- 60% file handle usage

Note: 60% is chosen to afford time and resources to add more nodes and move partitions accordingly

Kafka

Kafka publishes JMX metrics, documented here: <http://kafka.apache.org/documentation.html#monitoring>

The below table has additional information about some broker JMX metrics listed in the documentation:

Metric	Description
<code>kafka.server:type=BrokerTopicMetrics,name=MessagesInPerSec</code>	Number of incoming messages per second. Useful for understanding broker load.
<code>kafka.server:type=BrokerTopicMetrics,name=BytesInPerSec</code>	Bytes in per second. Useful for understanding broker load.
<code>kafka.network:type=RequestMetrics,name=RequestsPerSec,request={Produce/FetchConsumer/FetchFollower}</code>	Number of requests per second, for produce, consumer fetch, and replica follower fetch. Useful for understanding broker load.
<code>kafka.server:type=BrokerTopicMetrics,name=BytesOutPerSec</code>	Bytes out per second. Useful for understanding broker load.
<code>kafka.log:type=LogFlushStats,name=LogFlushRateAndTimeMs</code>	The frequency of Kafka log page cache flushes, and how long the flushes take. Can be used to diagnose a slow disk or misconfigured page cache.
<code>kafka.server:type=ReplicaManager,name=UnderReplicatedPartitions</code>	Should always be 0. If > 0, likely a broker failed. Should alert when not 0.

Recently viewed articles

- [Confluent Platform Ops Checklist](#)
- [Monitoring Connect](#)
- [Monitoring Zookeeper](#)
- [ZooKeeper Rolling upgrade](#)
- [Kafka Broker Performance Diagnostics](#)

Related articles

- [Kafka Broker Performance Diagnostics](#)
- [Confluent Platform Ops Checklist](#)
- [Monitoring Zookeeper](#)
- [Monitoring Connect](#)
- [Kafka Broker runs out of memory due to maximum mmap count exceeded](#)

<i>kafka.controller:type=KafkaController,name=ActiveControllerCount</i>	Indicates if this broker is the controller. Should alert if this value changes because likely a broker failed.
<i>kafka.controller:type=ControllerStats,name=LeaderElectionRateAndTimeMs</i>	Rate and time of leader election. If non-zero, either a broker has failed, or one or many brokers are experiencing soft failures, causing thrashing.
<i>kafka.server:type=ReplicaManager,name=PartitionCount</i>	Number of replicas on this host. Shouldn't exceed 2-3k.
<i>kafka.server:type=ReplicaManager,name=IsrShrinksPerSec</i>	The rate at which the ISR is shrinking. Like LeaderElectionRateAndTimeMs, the ISR will shrink if a broker is shutdown, either gracefully or not. If the ISR is thrashing, likely brokers are soft failing.
<i>kafka.network:type=RequestMetrics,name=TotalTimeMs,request={Produce/FetchConsumer/FetchFollower}</i>	Total time a request takes to be completed, for produce, consumer fetch, and replica follower fetch. Where request time is spent can be understood by exploring the below metrics.
<i>kafka.network:type=RequestMetrics,name=RequestQueueTimeMs,request={Produce/FetchConsumer/FetchFollower}</i>	The time the request is waiting in the request queue, for produce, consumer fetch, and replica follower fetch. A high value can imply there aren't enough IO threads or the CPU is a bottleneck.
<i>kafka.network:type=RequestMetrics,name=ResponseQueueTimeMs,request={Produce/FetchConsumer/FetchFollower}</i>	The time the request is waiting in the response queue, for produce, consumer fetch, and replica follower fetch. A high value can imply there aren't enough network threads.
<i>kafka.network:type=RequestMetrics,name=LocalTimeMs,request={Produce/FetchConsumer/FetchFollower}</i>	The time the request is being processed by the leader locally, for produce, consumer fetch, and replica follower fetch. Often a high value implies a slow disk.
<i>kafka.network:type=RequestMetrics,name=RemoteTimeMs,request={Produce/FetchConsumer/FetchFollower}</i>	The time the request is waiting on a remote client, for produce, consumer fetch, and replica follower fetch. A high value can imply a slow network connection. For fetch request, if the remote time is high, it could be that there is not enough data to give in a fetch response. This can happen when the consumer or replica is caught up and there is no new incoming data. If this is the case, remote time will be close to the max wait time, which is normal. Max wait time is configured via replica.fetch.wait.max.ms and fetch.max.wait.ms.
<i>kafka.network:type=RequestMetrics,name=ResponseSendTimeMs,request={Produce/FetchConsumer/FetchFollower}</i>	The time the request is being sent back to the client, for produce, consumer fetch, and replica follower fetch. A high value can imply there aren't enough network threads or the CPU or network is a bottleneck.
<i>kafka.network:type=SocketServer,name=NetworkProcessorAvgIdlePercent</i>	The average fraction of time the network threads are idle. Useful when paired with the above request-related JMX metrics.
<i>kafka.server:type=KafkaRequestHandlerPool,name=RequestHandlerAvgIdlePercent</i>	The average fraction of time the I/O threads are idle. Useful when paired with the above request-related JMX metrics.

	related JVM metrics.
<i>kafka.controller:type=KafkaController,name=ControllerState</i>	The state the controller is in, i.e. the event that is currently being processed. Some actions like partition reassignment may take a while and include many events (potentially interleaved with other events), but that doesn't change the fact that at most one event is processed at a time.
<i>kafka.controller:type=ControllerStats,name=ControllerChangeRateAndTimeMs</i>	rate and latency for the controller change state
<i>kafka.controller:type=ControllerStats,name=TopicChangeRateAndTimeMs</i>	rate and latency for the controller to create new topics
<i>kafka.controller:type=ControllerStats,name=TopicDeletionRateAndTimeMs</i>	rate and latency for the controller to delete topics
<i>kafka.controller:type=ControllerStats,name=PartitionReassignmentRateAndTimeMs</i>	rate and latency for the controller to reassign partitions
<i>kafka.controller:type=ControllerStats,name=AutoLeaderBalanceRateAndTimeMs</i>	rate and latency for the controller to auto balance the leaders
<i>kafka.controller:type=ControllerStats,name=ManualLeaderBalanceRateAndTimeMs</i>	rate and latency for the controller to manually balance the leaders
<i>kafka.controller:type=ControllerStats,name=IsrChangeRateAndTimeMs</i>	rate and latency for the controller to manually balance the leaders
<i>kafka.controller:type=ControllerChannelManager,name=TotalQueueSize</i>	total size of the queue in ControllerChannelManager
<i>kafka.controller:type=ControllerChannelManager,name=QueueSize,brokerId=10</i>	QueueSize in ControllerChannelManager for a particular broker
<i>kafka.cluster:type=Partition,name=FailedIsrUpdatesPerSec</i>	measure the occurrences of failed ISR update in ZK
<i>kafka.server:name=UnderMinIsrPartitionCount,type=ReplicaManager</i>	The value of this metric is the total number of leader partitions on this broker whose in-sync replicas < minIsr
<i>kafka.cluster:name=UnderMinIsr,type=Partition,topic={topic},partition={partition}</i>	The value of this metric is 1 if the broker is leader of this partition AND the number of in-sync replicas of this partition < minIsr of this partition. Otherwise it is 0.
<i>kafka.controller:type=KafkaController,name=GlobalTopicCount</i>	count of all topics within the cluster
<i>kafka.controller:type=KafkaController,name=GlobalPartitionCount</i>	count of partitions across all topics in the cluster
<i>kafka.network:type=RequestMetrics,name=ErrorsPerSec,request=apikey_name,error=error_code_name_</i>	number of errors per second
<i>kafka.server:type=BrokerTopicMetrics,name=FetchMessageConversionsPerSec,topic={[-.\\w]+}</i>	number of Fetch messages converted per second for a given topic
<i>kafka.server:type=BrokerTopicMetrics,name=ProduceMessageConversionsPerSec,topic={[-.\\w]+}</i>	number of Produce messages converted per second for a given topic

<i>kafka.network:type=RequestMetrics,name=MessageConversionsTimeMs,request={Produce or Fetch}</i>	Time in milliseconds spent on message format conversions.
<i>kafka.network:type=RequestMetrics,name=RequestBytes,request=apiKey</i>	Size of requests for each request type.
<i>kafka.network:type=RequestMetrics,name=TemporaryMemoryBytes,request=apiKey</i>	Temporary memory used for message format conversions and decompression.
<i>kafka.server:type=socket-server-metrics,listener=listenerName,networkProcessor=processorIndex</i>	Rate of successful/failed authentications using SASL or SSL
<i>kafka.server:type=ZooKeeperClientMetrics,name=ZooKeeperRequestLatencyMs</i>	Latency in milliseconds for ZooKeeper requests from broker.
<i>kafka.server:type=SessionExpireListener,name=SessionState</i>	Connection status of broker's ZooKeeper session which may be one of Disconnected
<i>_lkafka.admin.client</i>	kafka.consumer

JVM

JVM garbage collector (GC) logging should be enabled by setting the `GCLOG_ENABLED_` environment variable to `true`. The GC log will show how long garbage collection takes.

ZooKeeper

ZooKeeper can be monitored in the following ways:

- JMX, documented here: <https://zookeeper.apache.org/doc/r3.3.3/zookeeperJMX.html>
- Four-letter words, documented here: https://zookeeper.apache.org/doc/r3.3.3/zookeeperAdmin.html#sc_zkCommands
 - In particular, RUOK should return IMOK.
- Monitor the request latency and number of outstanding requests, via the following JMX metrics (created by the ZooKeeper server):
 - On a standalone ZooKeeper server:
 - `org.apache.ZooKeeperService:type=StandaloneServerport{port-number} AvgRequestLatency_`
 - `org.apache.ZooKeeperService:type=StandaloneServerport{port-number} OutstandingRequests_`
 - On a clustered, multi-machine ZooKeeper server:
 - `org.apache.ZooKeeperService:name0=ReplicatedServerid{myid},name1=replica.{myid},name2=Follower AvgRequestLatency_`
 - `org.apache.ZooKeeperService:name0=ReplicatedServerid{myid},name1=replica.{myid},name2=Follower OutstandingRequests_`
- Monitor which ZooKeeper servers are leaders and followers – this is useful when diagnosing issues. This information is published via JMX from the ZooKeeper server.
- Monitor the number of clients and watchers, via the following JMX metrics (created by the ZooKeeper server):
 - On a standalone ZooKeeper server:
 - `org.apache.ZooKeeperService:type=StandaloneServerport{port-number} NumAliveConnections_`
 - `org.apache.ZooKeeperService:type=InMemoryDataTree WatchCount`
 - On a clustered, multi-machine ZooKeeper server:
 - `org.apache.ZooKeeperService:name0=ReplicatedServerid{myid},name1=replica.{myid},name2=Follower NumAliveConnections_`
 - `org.apache.ZooKeeperService:name0=ReplicatedServerid{myid},name1=replica.{myid},name2=Follower,name3=InMemoryDataTree WatchCount_`

Consumer Lag

For "real-time" consumer applications, where the consumer is meant to be processing the newest messages with as little latency as possible, consumer lag should be monitored closely. Most "real-time" applications will want little-to-no consumer lag, because lag introduces end-to-end latency.

Consumer lag can be monitored using the *kafka-consumer-groups* command-line tool, or using the consumer's JMX metric *kafka.consumer:type=consumer-fetch-manager-metrics,client-id={client-id} attribute: records-lag-max*.

The *kafka-consumer-groups* command-line tool calculates lag by comparing the last committed offset to the most recent offset, which means lag is a function of when the last commit was made and hence may not be up-to-date. On the other hand, the consumer's records-lag-max JMX metric calculates lag by comparing the offset most recently seen by the consumer to the most recent offset in the log, which is a more real-time measurement.

Was this article helpful?

3 out of 3 found this helpful

f

t

in

Have more questions? [Submit a request](#)

1 Comments

Date

Votes

0

Sridhar Shanmugham

January 22, 2019 15:52

Is there any JMX metrics to monitor the peak number of messages processed per topic per day ?

⚙

▼